

Homework Set 2: Local Feature Matching

Instructions

- This homework contains two parts, i.e., a set of questions (20 pts, in Section 1) and a programming project (80 pts, in Section 2). The programming project has extra credit (maximum 10 pts).
- Upload the codes and electronic version of your report in a zipped file in the form of “`id_yourname.zip`” via this [link](#). The report includes the answers of the questions in Section 1; and describes the algorithm process, shows the results, describe any extra credit in Section 2 and 3.
- The deadline is May 19th, 2021.

1 Questions (20 pts)

1.1 Imagine we wished to find points in one image which matched to the same world point in another image – so-called feature point correspondence matching. We are tasked with designing an image feature point algorithm which could match world points in the following three pairs of images:



Figure 1: These three pairs of images are in the directory `data/matching`.

Please use the included python script `plot_corners.py` to find corners using Harris corner detection for each image above.

For each pair, discuss the differences in the returned corners (if any), and what aspects of the images may have caused these differences. Then discuss what real world challenges exist in matching features between images.

1.2 The Harris Corner Detector is commonly used in computer vision algorithms to find locations from which to extract stable features for image matches.

How do the eigenvalues of the \mathbf{M} second moment matrix vary with local image brightness, and how might we interpret the eigenvalues geometrically?

1.3 Given a feature point location, the SIFT algorithm converts a 16×16 patch around the feature point into a 128×1 descriptor of the gradient magnitudes and orientations therein. Write pseudocode with *matrix/array indices* for these steps.

Notes: Do this for just one feature point at one scale; ignore the overall feature point orientation; ignore the Gaussian weighting; ignore all normalization post-processing; ignore image boundaries; ignore sub-pixel interpolation and just pick an arbitrary center within the 16×16 for your descriptor. Please just explain in pseudocode how to go from the 16×16 patch to the 128×1 vector. You are free to simplify the gradient computation.

```
# You can assume access to the image, x and y gradients, and their magnitudes/
# orientations.

image = imread('rara.jpg')
grad_x = filter(image, 'sobelX')
grad_y = filter(image, 'sobely')
grad_mag = sqrt( grad_x.^2 + grad_y.^2 )
grad_ori = atan2( grad_y, grad_x )

# Takes in a feature point x,y location and returns a descriptor
def SIFTdescriptor(x, y)
    descriptor = zeros(128,1)

    # write your psudocode here ...

    return descriptor
```

1.4 Distance metrics for feature matching.

- (a) Explain the differences between the geometric interpretations of the Euclidean distance and the cosine similarity metrics. What does this tell us about when each should be used?
- (b) Given a distance metric, what is a good method for feature descriptor matching and why?

1.5 Given a point $\mathbf{x} = (u, v)$ in the image plane, and a vector $\mathbf{n} = (a, b)$, any point \mathbf{y} following the equation $\mathbf{y} = \lambda\mathbf{n} + \mathbf{x}, \forall \lambda \in \mathbf{R}$, defines a line ℓ . Now given a homography matrix

$$\mathbf{H} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix}, \quad (1)$$

- (a) Calculate the new line ℓ' after transforming the line ℓ by the homography \mathbf{H} , i.e. how \mathbf{y}' will be after transforming from \mathbf{y} . Note that \mathbf{y}' is also in the image plane.
- (b) Discuss in which case the line ℓ' will converge to a vanishing point, and calculate the vanishing point \mathbf{y}'_v .
- (c) Discuss in which case the line ℓ' will not converge to a vanishing point.

2 Programming Project (80 pts)

2.1 Overview

We will create a local feature matching algorithm and attempt to match multiple views of real-world scenes. There are hundreds of papers in the computer vision literature addressing each stage. We will implement the Harris corner detection and a simplified version of [SIFT](#); however, you are encouraged to experiment with more sophisticated algorithms and the later on image alignment for extra credit!

Task: Implement the three major steps of local feature matching:

- **Detection** in the `get_interest_points` function in `student.py`. Please implement the Harris corner detector. You do not need to worry about scale invariance or keypoint orientation for your Harris corner detector.
- **Description** in the `get_features` function, also in `student.py`. Please implement a *SIFT-like* local feature descriptor. *You do not need to implement full SIFT!* Add complexity until you meet the rubric. To quickly test and debug your matching pipeline, start with normalized patches as your descriptor.
- **Matching** in the `match_features` function of `student.py`. Please implement the “ratio test” or “nearest neighbor distance ratio test” method of matching local features.

Potentially useful functions: Any existing filter function, `zip()`, `skimage.measure.regionprops()`, `skimage.feature.peak_local_max()`, `numpy.arctan2()`.

Potentially useful libraries: `skimage.filters.x()` or `scipy.ndimage.filters.x()`, which provide many pre-written image filtering functions. `numpy.gradient()`, which provides a more sophisticated estimate of derivatives. `numpy.digitize()` provides element-wise binning. In general, anything which we’ve implemented ourselves in a previous project is fair game to use as an existing function.

Forbidden functions: `skimage.feature.daisy()`, `skimage.feature.ORB()`, and any other functions that extract features for you, `skimage.feature.corner_harris()` and any other functions that detect corners for you, any function which computes histograms, `sklearn.neighbors.NearestNeighbors()` and any other functions that compute nearest neighbor ratios for you. `scipy.spatial.distance.cdist()` and any other functions that compute the distance between arrays of vectors (use the guide we’ve provided to implement your own distance function!). If you are unsure about a function, please ask.

2.2 Running the code

`main.py` takes a command-line argument using ‘-p’ to load a dataset, e.g., ‘-p notre_dame’. For example, ‘\$ python main.py -p notre_dame’. Please see `main.py` for more instructions.

2.3 Requirements/Rubric

If your implementation reaches 60% accuracy on the *most confident* 50 correspondences in ‘matches’ for the Notre Dame pair, and at least 60% accuracy on the *most confident* 50 correspondences in ‘matches’ for the Mt. Rushmore pair, you will receive 70 pts (full code credit). We will evaluate your code on the image pairs at `scale_factor=0.5` (`main.py`), so please be aware if you change this value. The evaluation function we will use is `evaluate_correspondence()` in `helpers.py`. We have included this function in the starter code for you so you can measure your own accuracy.

Time limit: The time limit of executing your code is 20 minutes, after which you will receive a maximum of 40 pts for the implementation. You must write efficient code.

- Hint 1: Use ‘steerable’ (oriented) filters to build the descriptor.
- Hint 2: Use matrix multiplication for feature descriptor matching.

Compute/memory limit: If you are concerned about the memory usage of your program, you can check it by running `python memusecheck` in your code directory.

- (25 pts) Implementation of Harris corner detector in `get_interest_points()`.
- (25 pts) Implementation of SIFT-like local feature in `get_features()`.

- (c) **(20 pts)** Implementation of “Ratio Test” matching in `match_features()`.
- (d) **(10 pts)** Report:
- (i) Please describe your process and algorithm, show your results, describe any extra credit, and tell us any other information you feel is relevant.
 - (ii) Quantitatively compare the impact of the method you’ve implemented. For example, using SIFT-like descriptors instead of normalized patches increased our performance from 50% good matches to 70% good matches. Please include the performance improvement for any extra credit.
 - (iii) Show how well your method works on the Notre Dame, Mount Rushmore, and Episcopal Gaudi image pairs — the visualization code will write image-pair-specific output to your working directory, e.g. `notre_dame_matches.jpg`
- (e) **(10 pts)** Extra credit (max 10 pts total, thus there is no need to solve all questions).
- (i) Detection (**up to 5 pts**): Detect keypoints at multiple scales or use a scale-selection method to pick the best scale.
 - (ii) Description:
 - 1) **(Up to 3 pts)**: Experiment with SIFT parameters: window size, number of local cells, orientation bins, different normalization schemes.
 - 2) **(Up to 5 pts)**: Estimate feature orientation.
 - 3) **(Up to 5 pts)**: Multi-scale descriptor. If you are detecting keypoints at multiple scales, you should build the features at the corresponding scales, too.
 - (iii) Matching. The baseline matching algorithm is computationally expensive, and will likely be the slowest part of your code. Consider approximating or accelerating feature matching:
 - 1) **(Up to 5 pts)**: Create a lower dimensional descriptor that is still sufficiently accurate. For example, if the descriptor is 32 dimensions instead of 128 then the distance computation should be about 4 times faster. PCA would be a good way to create a low dimensional descriptor. You would need to compute the PCA basis on a sample of your local descriptors from many images.
 - 2) **(Up to 5 pts)**: Use a space partitioning data structure like a KD-tree or some third party approximate nearest neighbor package to accelerate matching.
 - 3) **(Up to 5 pts)**: Remove outlier matches if exist. You may use RANSAC together with homography estimation, and then report the output homography matrix for each test image pair.

2.4 An Implementation Strategy

1. Use `cheat_interest_points()` instead of `get_interest_points()`. At this point, just work with the Notre Dame image pair (the cheat points for Mt. Rushmore aren’t very good and the Episcopal Palace is difficult to feature match). This function CANNOT be used in your final implementation. It directly loads the 100 to 150 ground truth correspondences for the test cases. Even with this cheating, your accuracy will initially be near zero because the features are random and the matches are random.
2. Implement `match_features()`. Accuracy should still be near zero because the features are random. To do this, you’ll need to calculate the distance between all pairs of features across the two images (much like `scipy.spatial.distance.cdist()`). While this could be written using a series of `for` loops, we’re asking you to write a matrix implementation using `numpy`. Make sure to also return proper confidence values for each match. Your most confident matches should be highest values.

3. Change `get_features()` to cut out image patches. Accuracy should increase to $\sim 60 - 70\%$ on the Notre Dame pair if you're using 16×16 (256 dimensional) patches as your feature. Accuracy on the other test cases will be lower (especially using the bad interest points from `cheat_interest_points()`). Image patches are not invariant to brightness change, contrast change, or small spatial shifts, but this provides a baseline.
4. Finish `get_features()` by implementing a SIFT-like feature. Accuracy should increase to 70% on the Notre Dame pair. If your accuracy doesn't increase (or even decreases a little bit), don't worry because normalized patches are actually a good feature descriptor for the Notre Dame image pair. The difference will become apparent when you implement `get_interest_points()` and can test on the Mt. Rushmore pair.
5. Stop cheating and implement `get_interest_points()`. Harris corners aren't as good as ground-truth points (which we know to correspond), so accuracy may drop. On the other hand, you can get hundreds or even a few thousand interest points, so you have more opportunities to find confident matches. *Our solution* generates around $300 \sim 700$ feature points for each image.

These feature point and accuracy numbers are only a guide; don't worry if your method doesn't exactly produce these numbers. Feel confident if they are approximately similar, and move on to the next part.

2.5 Notes

`main.py` handles files, visualization, and evaluation, and calls placeholders of the three functions to implement. For the most part you will only be working in `student.py`.

The Notre Dame, Mount Rushmore, and Episcopal Gaudi image pairs include ‘ground truth’ evaluation. `evaluate_correspondence()` will classify each match as correct or incorrect based on hand-provided matches. You can test on those images by changing the `-p` argument you pass to `main.py`.

As you implement your feature matching pipeline, check whether your performance increases using `evaluate_correspondence()`. Take care not to tweak parameters specifically for the initial Notre Dame image pair.

You will likely need to do extra credit to get high accuracy on Episcopal Gaudi.

Feedback? (Optional)

Please help us make the course better. If you have any feedback for this assignment, we'd love to hear it!